

NIRAV SHAH

MOUNTAIN VIEW, CA — (984) 683-2844 — niravsshaha@gmail.com — linkedin.com/in/niravsshaha2

SUMMARY

Data Engineer with 4+ years of experience building batch and streaming pipelines on GCP and Azure. At Yahoo, cut core pipeline cloud cost 94% (saving \$150K+/month), maintained 50+ pipelines processing 5B+ events/day at 99.9% SLA, and led an 8-engineer migration to Cloud Composer. Promoted to Data Engineer II within 2 years.

EXPERIENCE

Yahoo (Mountain View, CA)

Data Engineer II

Mar 2026 - Present

- Led migration of 5 core event pipelines from Pig on Dataproc to 3 partitioned BigQuery SQL jobs, cutting cloud cost 94% (\$161K to \$9K/month) and runtime 85%; authored the design doc reused for later migrations.
- Built a near real-time GA4 streaming pipeline on BigQuery and Airflow processing 4M views/hour at 10-minute freshness for \$1.2K/month, used by 100+ dashboard users across 10+ Yahoo teams.
- Diagnosed a data quality defect inflating engagement rate 28% for Yahoo News/Homepage and Yahoo Sports, tracing 884K phantom sessions to GA4 event logic and filtering them out in the content ETL.
- Automated schema changes (new columns and metrics) with a Claude AI agent workflow, cutting each change from 2 weeks to 2 days and shipping 20+ metrics, including users, impressions, and sessions.
- Restored content attributes for 85% of real-time hourly views after an upstream source was retired; built, deployed, and validated a Apache Beam Dataflow fix that permanently moved enrichment to the new source.

Data Engineer I

Feb 2024 - Feb 2026

- Led a team of 8 engineers migrating batch workflows from on-prem Oozie and Kubernetes to GCP Cloud Composer, redesigning Oozie ETL workflows to scalable Airflow DAGs and cutting peak runtime 20%.
- Maintained 50+ batch and streaming pipelines (Pub/Sub, Kafka, Storm; 5B+ events/day) at 99.9% SLA as on-call owner, integrating OpsGenie alerting with Airflow and cutting incident response time 50%.
- Engineered BigQuery ML (ARIMA) anomaly detection across 30 content and video KPIs, catching 10+ anomalies in 3 months, including a homepage URL misroute fixed in 5 hours; design doc adopted by 2 teams.
- Redesigned referral classification logic for a certified audience-measurement pipeline, reducing uncategorized pageviews 92% (20M to 1.5M/day) and improving downstream reporting accuracy.

Data Engineer Intern

May 2023 - Aug 2023

- Built a BigQuery query benchmarking tool by extending a REST API to tag BigQuery jobs, testing user queries behind 7+ API endpoints and flagging 50 slow, resource-intensive queries for optimization.
- Designed the benchmark data flow: tagged each BigQuery job with a UUID, date window, and range label, joined runs with INFORMATION_SCHEMA for runtime and resource usage, and stored results in Spanner.

Honeywell (Pune, India)

Data Engineer

Oct 2020 - May 2022

- Engineered Spark/Databricks batch pipelines loading 100+ chemical plant params from PI Historian to Azure Data Lake at 70% lower latency, feeding Tableau and Power BI dashboards that cut resolution time 40%.
- Developed AutoML and time-series ML models on process parameters that recommended optimal plant settings and predicted equipment failures, contributing to a 5% manufacturing throughput gain.

EDUCATION

North Carolina State University (Master of Computer Science) GPA: 4.0/4.0 Aug 2022 - Dec 2023

Pune Institute of Computer Technology (B.E. in Computer) GPA: 8.99/10 Aug 2016 - May 2020

TECHNICAL SKILLS

Languages: Python, SQL, Java, C, C++, JavaScript
Cloud: GCP (BigQuery, Dataflow, Composer, Spanner), AWS, Azure (Databricks, Data Lake)
Data Engineering: Airflow, ETL/ELT, Pub/Sub, Kafka, Storm, Apache Spark, Hive, Hadoop
DevOps: CI/CD, Terraform, Docker, Kubernetes, Jenkins, GitHub Actions, Git, pytest
BI & Monitoring: Looker (LookML), Tableau, Power BI, Rootly, BigPanda, OpsGenie